

Distributing OpenBUGs chains in parallel

Tom Werner
Chris Chekal

Department of Computer Science
The University of Iowa
Iowa City, IA

December 2, 2016

What is OpenBUGs?

- ▶ Software for **B**ayesian inference **U**sing **G**ibbs **S**ampling
- ▶ Uses Monte Carlo Markov Chains (MCMC) to model posterior distributions

Why use OpenBUGs?

- ▶ Free and open source software
- ▶ Integrates with R through R2OpenBugs/Coda packages
- ▶ Used in Bayesian analysis class at Iowa

What is being parallelized?

- ▶ Using MCMC requires drawing many samples from a converged chain
- ▶ Determining convergence is easier with multiple chains
- ▶ If chains converge to different spots, something is wrong
- ▶ Each chain is independent of the others, allowing for naive parallelization
- ▶ Just need to make sure they're seeded independently

Easy cluster/CPU parallization in R

- ▶ We used the Snow and Snowfall packages
- ▶ They let you easily set up a cluster to run R code on
- ▶ sfLapply works like regular apply, but distributes it over the cluster

Snow/Snowfall Example

- ▶

```
data = matrix(rnorm(1000 * 1000), ncol=100)
library(snow)
library(snowfall)
sfInit(parallel=TRUE, cpus=4)
print(system.time(sfLapply(data, mean)))
sfStop()
sfInit(parallel=TRUE, cpus=1)
print(system.time(sfLapply(data, mean)))
sfStop()
```
- ▶ Parallel: 2.61 seconds
- ▶ Sequential: 6.26 seconds

Parallel Bugs workflow

1. Define openbugs model as a string, data as an R list, and inits as a list
2. Call parallelBugs with these arguments, and the IP addresses of cluster machines
3. Use Coda or other R packages for analyzing MCMC output

Parallel Bugs workflow

1. Define openbugs model as a string, data as an R list, and inits as a list

```
model.str = "model
{
  tau ~ dgamma(0.0001, 0.0001)
  mu ~ dflat()
  sigma <- 1 / sqrt(tau)

  for(i in 1:N)
  {
    x.data[i]~dnorm(mu, tau)
  }
}"
```


Parallel Bugs workflow

1. Define openbugs model as a string, data as an R list, and inits as a list

```
# Simulated normal dataset
bugs.data = list(x.data=rnorm(1000, 12, 5), N=1000)

# Initial values of parameters. Note that you can use
# R random number functions here
init.list = list(
  chain1=list(mu=runif(1, -50, 50), tau=1),
  chain2=list(mu=-50, tau=10),
  chain3=list(mu=50, tau=.1)
)

# Parameters to collect statistics on
params <- c("mu", "sigma")
```

Does it help?

1. Parallel time for 100,000 iterations, 3 chains: 60 seconds
2. Sequential time: 115 seconds

```
system.time(bugs(data=bugs.data, inits=init.list,  
                 parameters.to.save=params, n.iter=100000,  
                 n.chains=3, n.burnin=1000, n.thin=1,  
                 model.file="model.txt", debug=FALSE,  
                 codaPkg=TRUE, OpenBUGS.pgm=OpenBUGS.exe,  
                 working.directory=folder,  
                 bugs.seed=as.integer(runif(1, 1, 14))))
```

```
user    system elapsed  
0.17     0.03    60.69
```

```
system.time(parallel.bugs('localhost', 3, model.str,  
                           bugs.data, 3, init.list, params,  
                           OpenBUGS.exe, n.iter=100000))
```

```
user    system elapsed  
0.42     0.00   115.22
```

References

- ▶ Luke Tierney, A. J. Rossini, Na Li and H. Sevcikova (2016). snow: Simple Network of Workstations. R package version 0.4-2. <https://CRAN.R-project.org/package=snow>
- ▶ Jochen Knaus (2015). snowfall: Easier cluster computing (based on snow).. R package version 1.84-6.1. <https://CRAN.R-project.org/package=snowfall>
- ▶ Martyn Plummer, Nicky Best, Kate Cowles and Karen Vines (2006). CODA: Convergence Diagnosis and Output Analysis for MCMC, R News, vol 6, 7-11